

Citizen AI Project Documentation

1. Introduction

Project Title: Citizen AI – Intelligent Citizen Engagement Platform

Team Id: NM2025TMID00797

Team Members:

- Priya Dharshini P
- SelvaPriya S
- Pooja V

2. Project Overview

Purpose

Citizen AI is an advanced citizen engagement platform designed to transform the way governments connect with the public. Built using Flask, IBM Granite LLM, and IBM Watson, it delivers real-time, AI-powered responses to citizen questions about government services, policies, and civic matters.

By incorporating natural language processing (NLP) and sentiment analysis, the platform evaluates public sentiment, identifies emerging concerns, and provides actionable insights for government agencies.

Citizen AI enhances citizen satisfaction, streamlines government operations, and fosters greater public trust in digital governance by automating routine communications and supporting data-driven policy making.

Features

- Conversational Chat Assistant
 - o Highlight: Real-time interaction with citizens
 - o Description: Delivers immediate, human-like AI responses to inquiries about government services, policies, and civic concerns.
- Citizen Sentiment Analysis
 - o Highlight: Monitoring public sentiment
 - o Description: Analyzes citizen feedback and categorizes it as Positive, Neutral, or Negative to identify satisfaction levels and areas of concern.
- Dynamic Analytics Dashboard
 - o Highlight: Insight-driven decision-making
 - o Description: Presents visual data on trends, public sentiment, and reported issues to support informed policymaking.

- Concern Reporting
 - Highlight: Transparent issue resolution
 - Description: Enables citizens to submit complaints or concerns, which are tracked for follow-up and resolution.

Use Case Scenarios

1. Interactive AI Chat Assistant
Citizens interact through a chat interface, receiving prompt and accurate responses from the AI regarding public services, government policies, and civic matters.
2. Sentiment Analysis of Citizen Feedback
The system processes citizen input, evaluates sentiment (Positive, Neutral, or Negative), and compiles the data to support informed decision-making.
3. Real-Time Analytics Dashboard
Government officials access a live dashboard that displays trends in public sentiment, service satisfaction, and reported issues, allowing for timely and data-driven responses.

3. Architecture

- Frontend:
Built with HTML and CSS, the interface includes templates for the homepage, about section, services, chat interface, analytics dashboard, and user login.
- Backend (Flask):
Responsible for handling application routes, user authentication, and backend data logic.
- Large Language Model (IBM Granite):
Integrates advanced AI capabilities for natural conversation, text generation, and sentiment detection.
- Data Management:
Currently uses in-memory storage to manage chat logs, sentiment data, and reported concerns, with future plans to implement a database solution.
- Data Visualization:
Includes a dynamic dashboard that presents real-time charts and analytics for tracking sentiment patterns and issue reports.

4. Setup Instructions

Prerequisites

- o Python 3.7+
- o Flask
- o PyTorch (with CUDA for GPU acceleration)
- o Hugging Face libraries: transformers, accelerate, bitsandbytes
- o Hardware:
 - 16GB+ RAM
 - NVIDIA GPU with 8GB+ VRAM recommended
- o Internet connection (for first-time model download)

Installation Process

1. Clone repository and set up project structure (app.py, templates/, static/).
2. Create and activate a virtual environment:
3. `python -m venv env`
4. `source env/bin/activate` # Linux/Mac
5. `env\Scripts\activate` # Windows
6. `pip install -r requirements.txt`
7. Install Flask, PyTorch, and Hugging Face dependencies.
8. Configure IBM Granite model path (ibm-granite/granite-3.3-8b-instruct).
9. Run the Flask backend with:

10. `python app.py`

5. Folder Structure

app.py – Main Flask application

templates/ – HTML templates (index, about, services, chat, dashboard, login)

static/ – CSS, Images, Favicon

requirements.txt – Python dependencies

6. Running the Application

- o Launch Flask backend (`python app.py`).
- o Open browser and navigate to `http://localhost:5000`.
- o Use navigation menu for:

Chat: Interact with the AI assistant.

Feedback: Submit text for sentiment analysis.

Dashboard: View real-time citizen insights.

Login: Authenticate to access protected content.

7. API Endpoints

- POST /ask – Accepts citizen inquiries and returns AI-generated responses.
- POST /feedback – Submits user feedback for sentiment analysis.
- POST /concern – Allows users to report issues or concerns.
- GET /dashboard – Retrieves aggregated sentiment data and reported issues.
- POST /login – Handles user login and authentication.
- POST /logout – Ends the current user session.

8. Authentication

- Supports user login with session-based authentication.
- Upcoming Feature: Role-based access control for different user types (citizens, administrators, government officials).

9. User Interface

- Index Page: Introductory landing page with a welcome message and "Get Started" call to action.
- Login Page: Secure login form for user authentication.
- About Page: Describes the platform's mission, key features, and benefits to users.
- Chat Page: Interactive interface for AI-powered conversations with citizens.
- Dashboard: Displays sentiment breakdown (positive, neutral, negative) and recent citizen-reported issues in real time.

10. Testing Strategy

- Unit Testing: Covers Flask routes and core AI functionalities.
- Integration Testing: Validates complete workflows across chat, feedback submission, and dashboard updates.
- Manual Testing: Includes form submissions, sentiment classification accuracy, and issue reporting validation.
- Edge Case Handling: Tests for scenarios like empty input fields, invalid login attempts, and improperly formatted data.

11. Screenshots

hugging face

Hugging Face
https://huggingface.co/

Hugging Face – The AI community building the future.
The AI community building the future. The platform where the machine learning community collaborates on models, datasets, and applications. Explore AI Apps.

Kolours Virtual Try-On in the Wild
Upload a person image and a garment image to see how the ...

Spaces
Image Generation · Video Generation · Text Generation ...

Log in
Write on a journey to advance and democratize artificial ...

Learn
Write on a journey to advance and democratize artificial ...

Models
Explore machine learning models

Please check your email address for a confirmation link. [Resend confirmation email](#)

ibm-granite/granite-3.2-2b-instruct

Text Generation · Transformers · Safetensors · granite · language · granite-3.2 · conversational · arXiv:0000.00000

License: apache-2.0

Model card · Files · set · Community

Train · Deploy · Use this model

A newer version of this model is available: [ibm-granite/granite-3.3-2b-instruct](#)

Downloads last month
16,448

Safetensors

Granite-3.2-2B-Instruct

google colab

Welcome To Colab - Colab - Google

Colab notebooks allow you to combine executable code and rich text in a single document, along with images, HTML, LaTeX and more. When you create your own Colab ...

Jupyter.ipynb
This section describes how to edit and run the code in each ...

Run in Google Colab
Colab Specifics ... Colab is a virtual machine you can access ...

Pro
Colab Pro ... 1 limited time offer of an additional 400 compute hours

Citizen AI

File Edit View Insert Runtime Tools Help

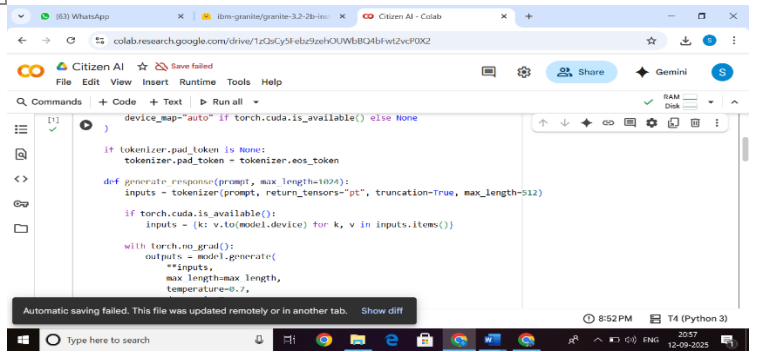
Commands + Code + Text + Run all

Connect 14

```
run this project file in google colab by changing run type to T4 GPU
!pip install transformers torch gradio -q
import gradio as gr
import torch
from transformers import AutoTokenizer, AutoModelForCausalLM

# Load model and tokenizer
model_name = "ibm-granite/granite-3.2-2b-instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
    device_map="auto" if torch.cuda.is_available() else None
)
```

Variables Terminal



/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret 'HF_TOKEN' does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret in:
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
tokenizer_config.json: 8.88k/? [00:00:00.00, 100kB/s]
vocab.json: 777k/? [00:00:00.00, 9.07MB/s]
merges.txt: 442k/? [00:00:00.00, 6.65MB/s]
tokenizer.json: 3.48M/? [00:00:00.00, 10.6MB/s]
added_tokens.json: 100% 87.0/87.0 [00:00:00.00, 2.67kB/s]
special_tokens_map.json: 100% 701/701 [00:00:00.00, 27.7kB/s]
config.json: 100% 786/786 [00:00:00.00, 21.2kB/s]
`torch_dtype` is deprecated! Use `dtype` instead!
model.safetensors.index.json: 29.8k/? [00:00:00.00, 2.19MB/s]
Fetching 2 files: 0% 0/2 [00:00:00.00, 74B/s]
model-00002-of-00002.safetensors: 100% 67.1M/67.1M [00:01:00.00, 69.3MB/s]
model-00001-of-00002.safetensors: 76% 3.79G/5.00G [01:15:00.24, 49.2MB/s]

12. Current Limitations

- Data is temporarily stored in memory (no long-term persistence).
- Performance is slow due to CPU-only execution.
- Scalability is restricted until database support is added.

13. Planned Enhancements

- Integration with a database for reliable data storage.
- Support for multiple languages.
- Fully responsive, mobile-friendly design.
- Adoption of advanced NLP models for improved policy summarization.
- Connectivity with social media and public forums for sentiment analysis.